



Universal tokenization metrics (mean $\pm$ std across 4 datasets $\times$ 3 seeds)

Compression ratio (higher is better)



Reconstruction MSE (lower is better)



ACF retention (higher is better)

